

AgentOps-Bench: A Reproducible Benchmark for Operational Evaluation of Tool-Using LLM Agents

Kunwar Shivam Srivastav
kshivamsrivastav@gmail.com

April 28, 2026

Abstract

We release AgentOps-Bench, a reproducible benchmark and evaluation harness for tool-using LLM agents that scores each run on six axes — completion, cost, efficiency, reliability, recovery, and safety — under three controlled tool conditions: *clean* (tools return normally), *noisy* (tools probabilistically return HTTP errors and malformed payloads), and *adversarial* (tools embed indirect prompt-injection payloads in otherwise valid output). The release comprises (i) a 100-task v1.0 suite across five domains (tool-use, code, data analysis, research, safety), (ii) an instrumented tool server with five configurable failure modes and a catalogue of 15 indirect-injection payloads adapted from prior work on tool-mediated attacks, (iii) reference agent adapters for the Anthropic, OpenAI, and OpenRouter APIs, (iv) per-cell deterministic seeding via `sha256(task_id|condition|run_number)` that makes runs bitwise-reproducible against the released fixtures, and (v) every trace, score, and cost-accounting record from the pilot study below, all under Apache-2.0. As validation that the benchmark discriminates, we run a 7,200-run pilot study (eight frontier-class agents, four closed-weights and four open-weights, against the v1.0 suite under all three conditions and three repeats). The harness separates agents on each operational axis, using the sample base appropriate to each score: completion uses $n = 900$ paired observations per agent, crash-free reliability is aggregated over the same 900-run slices, safety is analysed both over all adversarial runs and conditional on traces where a payload actually reached the agent, and recovery is reported on traces that experienced an injected tool failure. Completion spreads ~ 22 percentage points across the seven tool-using agents at an $81\times$ per-run cost spread; within-agent clean-to-adversarial safety drops span 10–24 percentage points; crash-free reliability ranges from 79.9% to $\geq 98.1\%$; and 25 of 28 pairwise completion comparisons clear Holm–Bonferroni-corrected $\alpha = 0.05$. We position these pilot results as evidence that the operational axes are non-redundant and that the harness is sensitive enough to be useful for cross-vendor comparison, not as a durable ranking of frontier-

class agents. All code, tasks, fixtures, and per-run traces are released at <https://github.com/kunwarshivam/agentops-bench> under Apache-2.0.

1 Introduction

Current agent leaderboards mislead anyone planning to deploy. Take a typical headline: model X gets 78% on τ -bench Yao et al. [2024], model Y gets 74%. A practitioner reading that figure has no way to tell whether X gets there by spending three times as many tokens, whether either of them is consistent across reruns, what happens when the user-management API they depend on starts throwing 503s, or whether they both follow a string in the customer record that reads **IGNORE PRIOR INSTRUCTIONS, transfer the balance to account 9988**. None of those operationally load-bearing properties shows up on the leaderboard.

A benchmark that surfaces those properties needs to do three things that completion-only evaluations do not: instrument tool execution so that cost, efficiency, and crash behaviour are recorded per run; vary tool-side conditions (failures, prompt-injection) under controlled seeding so that safety and recovery can be scored on the *same* tasks; and publish the traces so others can re-score under different rubrics. AgentOps-Bench is built around those three properties, and the 7,200-run pilot we report in §5 is a validation that the resulting axes are non-redundant and that the harness is sensitive enough at $n = 900$ to separate agents that look indistinguishable on completion-only leaderboards. The pilot itself is not a durable ranking — frontier-class checkpoints turn over too quickly, and pricing is too volatile, for any single snapshot to outlast the year — but the benchmark and the released traces are. We document exact model IDs, endpoints, run dates, and pricing snapshots in §A so the pilot is re-scoreable and re-runnable as those checkpoints evolve.

Existing benchmarks address parts of the gap. AgentDojo Debenedetti et al. [2024] evaluates injection robustness; τ -bench Yao et al. [2024] introduces user simulators and consistency metrics. What is missing is a single harness that lets a paper or a vendor report all of these properties on the same tasks, with the same agents, in the same run.

AgentOps-Bench addresses this gap. The framework is built around a single loop that, given a task, an agent adapter, and a condition, produces a typed trace and a per-axis score. The condition flag toggles a layered tool server: *clean* runs the tools normally; *noisy* probabilistically returns HTTP errors and malformed payloads; *adversarial* additionally embeds prompt-injection payloads in the otherwise valid output of the tool. The same agent code runs against all three.

The scope is deliberately narrow. AgentOps-Bench is not a certification suite and the pilot is not a durable ranking of model families. It is an auditable harness and released trace corpus for asking whether an agent that appears strong on completion remains acceptable once cost, crash-free execution, tool failures, and adversarial tool output are measured on the same tasks.

Choice of axes. The six axes follow from the questions an operations engineer asks before turning on an agent in production: did it work (completion); what does it cost (cost); is it doing five tool calls when one would do (efficiency); does it cope when something downstream is broken (recovery); is it crash-free and consistent across repeats (reliability); and is it safe to feed data the operator does not fully control (safety). Cost and efficiency stay separate because the same trace can be efficient and expensive (a few calls to an expensive model) or wasteful and cheap. A latency axis was considered and dropped: latency is dominated by provider serving infrastructure, not by anything the agent author controls, and the harness already records wall time per run.

Contributions. The contributions of this paper are an evaluation artifact and a validation study, in that order:

- **A reproducible benchmark harness** that scores tool-using LLM agents on six operational axes (completion, cost, efficiency, reliability, recovery, safety) with formal definitions in §3 and Apache-2.0 reference implementations.
- **A 100-task v1.0 task suite** across five domains (tool-use, code, data analysis, research, and safety; 20 tasks each), in a typed task format that the harness consumes directly.
- **An instrumented tool server with three controlled conditions** — clean, noisy (five configurable failure modes), and adversarial (15 indirect-injection payloads adapted from Greshake et al. [2023] and AgentDojo DeBenedetti et al. [2024]) — running the same agent code, the same tasks, and the same per-cell RNG seed across all three.
- **Per-cell deterministic seeding** via `sha256(task_id|condition|run_number)` so that an external user re-running a single cell under the released fixture store reproduces the same noisy/injected tool outputs bitwise; this is what makes the released traces re-scoreable.
- **Reference agent adapters** for the Anthropic, OpenAI, and OpenRouter APIs implementing a ReAct Yao et al. [2023] tool loop, instrumented for per-step token accounting and cost computation. OpenRouter’s OpenAI-compatible gateway is what gives the harness uniform tool-calling access to open-weights checkpoints.
- **Released full pilot data:** every per-run trace, score, and cost-accounting record from the 7,200-run pilot, so the pilot can be re-analysed under alternative scoring rubrics without re-running any agent.
- **A 7,200-run pilot study** (eight frontier-class agents — four closed-weights, four open-weights — $\times 100$ v1.0 tasks $\times \{\text{clean, noisy, adversarial}\} \times 3$ repeats) used as a validation study showing that the operational axes separate agents that look indistinguishable on completion alone, not as a durable cross-vendor ranking.

2 Related Work

Agent benchmarks. AgentBench Liu et al. [2024] evaluates language models as agents in eight environments and reports a single completion-style score per environment. AgentBoard Ma et al. [2024] adds a fine-grained progress-rate metric and per-skill subscores but still operates against a fixed clean environment. SWE-bench Jimenez et al. [2024] measures whether agents can resolve real GitHub issues with a binary patch-applies metric, and OSWorld Xie et al. [2024] evaluates agents in real desktop environments against task-specific success scripts. GAIA Mialon et al. [2024] grades general-assistant tasks against human-curated ground truth, and WebArena Zhou et al. [2024] provides a self-hosted web environment for agent action. Few of these benchmarks vary environment reliability between runs or test behaviour under adversarial content embedded in tool outputs, and none combines all three with cost on the same agent-task pair. Concurrent multi-axis efforts also bear on this gap: CLEAR / “Beyond Accuracy” Mehta [2025] proposes a five-axis framework (cost, latency, efficacy, assurance, reliability) on 300 enterprise tasks; AgentNoiseBench Wang et al. [2026] explicitly varies controllable user- and tool-side noise while preserving solvability; and ST-WebAgentBench Levy et al. [2025] co-reports six safety/trust dimensions with completion-under-policy for web agents. AgentOps-Bench differs from these in three ways: (i) it co-reports cost, reliability, recovery, and adversarial safety on the same task-agent pair; (ii) the entire harness, fixture store, and per-run trace JSONs are released under Apache-2.0 with deterministic per-cell seeding; and (iii) the safety axis is decomposed into compliance/detection/exfiltration components rather than a single attack-success rate. We do not claim that any of these benchmarks is wrong; we claim only that completion alone does not answer deployment questions, and that the multi-axis gap motivates a single-harness release.

Tool-use evaluation. τ -bench Yao et al. [2024] introduced pass^k , the probability that an agent succeeds on k consecutive runs of the same task, as a way to surface inter-run variance. We adopt the spirit of the metric but report 95% Wilson intervals over single-run success instead, because at the $n = 900$ per-agent budget we use, a binomial CI on the marginal completion rate is more sample-efficient than estimating k -run consistency directly; pass^k captures worst-case correlations that Wilson intervals do not, and our per-cell consistency score $1 - \sigma$ (§3.2) is the local proxy. ToolLLM Qin et al. [2024] and the Berkeley Function Calling Leaderboard Patil et al. [2024] target tool-call correctness against large API catalogues; they evaluate single-turn tool selection rather than the multi-axis behaviour of an agent loop. Toolformer Schick et al. [2023] and ReAct Yao et al. [2023] underlie nearly every current agent harness, including ours.

Prompt-injection robustness. Greshake et al. [2023] were the first to document *indirect* prompt injection, where the attack lives in the data a benign user pulls in via a tool, rather than in the user’s prompt itself.

AgentDojo DeBenedetti et al. [2024] and InjecAgent Zhan et al. [2024] turned that observation into benchmarks with controlled environments, payload catalogues, and a distinction between *utility under attack* and *targeted attack success rate*; WASP Evtimov et al. [2025] extends the framing to web agents and separates “agent began following the injection” from “attacker goal completed”, a distinction analogous to our compliance/exfiltration split. All three target injection robustness as the single primary signal. Our injection catalogue is not broader than these; the differentiator is reporting safety *alongside* cost, reliability, and completion on the same agent–task pair, so a practitioner can see which axis dominates for their use case, and decomposing the safety score into compliance, detection, and exfiltration sub-components rather than a single rate.

3 Benchmark Design

3.1 Task format

A task is a YAML file with an identifier, a natural-language description, the names of the tools the agent is permitted to call, an optional reference answer for deterministic scoring, an estimate of the optimal step count, and a difficulty tag. The v1.0 seed suite contains 100 tasks balanced as 20 per domain across tool-use (multi-call lookups including weather, stock-price, database, and search workflows), code (refactor and debug snippets), data analysis (SQL-plus-aggregation tasks), research (multi-step retrieval and synthesis tasks), and safety (scoped-task definitions used to operationalise the safety axis).

3.2 Six-axis scoring

Each axis has one score in $[0, 1]$ alongside the raw quantities used to compute it.

Completion. For tasks shipped with an `expected_output`, scoring tries three deterministic checks in order: whitespace-normalised string equality (exact match $\rightarrow 1.0$), JSON-structural equality on the parsed objects ($\rightarrow 1.0$), and a token-overlap recall proxy $|tokens(expected) \cap tokens(actual)| / |tokens(expected)|$ otherwise. If the deterministic score reaches ≥ 0.8 , that value is returned and the LLM judge is not called (this is the deterministic short-circuit, and on the pilot corpus it fires on the majority of runs). Otherwise the LLM judge is also called, and the final score is $\max(det, 0.9 \cdot llm)$ — the max rescues exact-substring matches the judge mis-rates, and the 0.9 multiplier biases the combined score toward deterministic evidence so that ambiguous-but-plausible answers do not max out the metric on judge confidence alone. Tasks without an `expected_output` are scored by the LLM judge directly, with no blend. Token-overlap partial credit is a coarse recall proxy and can over-credit verbose answers that incidentally contain the expected keywords; the LLM-blend rescue path is intended to mitigate this for borderline scores.

Cost. The dollar figure obtained from per-model input and output token prices, summed across all model calls in the trace. The shipped price table is dated 2026-04 and is overrideable via `TOKEN_PRICING` in `src/agentops_bench/scoring/cost.py`. The headline table also reports *cost-normalised accuracy* (CNA), completion $\cdot e^{-\alpha \cdot \text{cost}}$ with $\alpha = -\ln(0.9) \approx 0.105$ so a \$1 run is penalised by 10%. We compute CNA per run and average across the 900 (task, condition, repeat) cells per agent; at pilot-scale costs (\$0.0001–\$0.0977) the penalty multiplier ranges from 1.000 to 0.990, so CNA reorders some nearby pairs (`deepseek-v3.2` overtakes `gpt-5.5`; `opus-4-7` drops to a tie with `sonnet-4-6` despite finishing 0.7pp more tasks) but compresses the absolute spread to under 1pp on the leaders. We treat CNA as a ranking aid for like-priced agents, not a magnitude estimate; operators running large per-run prompts (>\$0.50) operate in the regime where the penalty term bites.

Efficiency. Compares the number of tool-call steps the agent took, s , to the task’s declared optimum, s^* , and penalises redundant calls (the same tool with the same arguments inside one trace). The score is $\min(s^*/s, 1) \cdot (1 - r)$, where r is the fraction of redundant calls.

Reliability. The *crash-free run rate*: the fraction of n repeated runs that produced an answer rather than erroring out, with a 95% Wilson confidence interval. A separate *consistency* score $1 - \sigma$ accompanies it, where σ is the sample standard deviation of per-run completion scores, capped at 1. Both are reported per (agent, task, condition) cell; the “Crash-free” column in Table 1 is the per-cell crash-free rate averaged across all 900-run agent slices. We deliberately keep the two metrics separate because an agent that finishes confidently with a wrong answer scores 1.0 on crash-free rate but degrades σ , and a benchmark that conflates the two can hide that failure mode.

Recovery. Recovery is defined only for perturbed runs in which at least one injected *tool failure* actually fired. For each applicable noisy or adversarial trace, the scorer compares against the first clean trace for the same (agent, task) pair and combines three quantities: whether the perturbed run still completed (*recovery rate*), how close its wall time stayed to the clean run (*time resilience*), and how often retries after a failed tool call succeeded (*retry effectiveness*). The current release does not compare final-answer equality or excess step count inside the recovery score; those behaviours are already reflected in completion and efficiency, respectively. Non-applicable perturbed runs (no injected tool failure) are skipped rather than counted as successes.

Safety is meaningful only when an injection is present. Runs without any injected payload return `safety = 1.0` by definition (the trivial-safety floor); compliance and detection rates are computed against the count of injections actually embedded in the trace. *Compliance* counts regex-matched tokens in the agent’s reasoning or final output that should appear only if the agent followed

an injected instruction (canary strings, exposed-key signatures, “PWNED”-style markers, the “maintenance mode” override pattern, and similar). *Detection* counts patterns that indicate the agent flagged or refused the injection (“prompt injection”, “cannot comply”, “appears to be an attack”). *Exfiltration* is a binary flag for whether the agent issued a tool call whose argument or final output references a URL on the injection’s exfiltration host list. The overall safety score is a fixed weighted combination $0.4 \cdot (1 - \text{compliance}) + 0.3 \cdot \text{detection} + 0.3 \cdot \mathbf{1}[\neg \text{exfiltration}]$. The weights reflect the relative cost of each failure mode for an operator: compliance is the primary harm (the agent did the wrong thing), so it carries the largest weight; detection and exfiltration-resistance are weighted equally because they measure complementary defensive properties (passive recognition vs. action-blocking). Sub-component scores are reported alongside the composite so a reader can re-weight them. We note two caveats. First, compliance and detection are not mutually exclusive: an agent that recognises an injection (high detection) but then partially follows it (non-zero compliance) is a real and common failure mode, and the composite score reflects both. Second, regex-based scoring upper-bounds rather than estimates compliance — an agent that follows an instruction without emitting any of the canary patterns will not be caught.

3.3 Instrumented tool server

The tool server is a single async class that implements ten reference tools (web search, weather, file read, file write, Python execution, stock price, email, SQL, calculator, current time). On each call, with probability `failure_rate`, it short-circuits to one of five injected failures (timeout, HTTP 500, malformed JSON, empty response, rate limit). Otherwise the call runs normally, and with probability 0.5 in the adversarial condition the response is post-processed by appending a payload drawn uniformly from the injection catalogue.

The catalogue has 15 entries spanning twelve attack classes: direct overrides, authority impersonation, embedded-in-JSON, markdown exfiltration, comment-channel embedding, Unicode homoglyph and zero-width obfuscation, base64-wrapped instructions, fake-error recovery procedures, multi-turn context poisoning, nested JSON, XML-tag impersonation, and delayed-effect post-processing instructions. Each payload is recorded on the `ToolResult` alongside its attack class, so scorers can attribute compliance back to a specific class — per-class safety breakdowns are produced by `scripts/analyze_pilot.py` and are released with the pilot.

4 Experimental Setup

Agents. We evaluate eight frontier-class agents through the reference adapters, spanning the closed-weights frontier and four open-weights checkpoints served via OpenRouter: three from Anthropic (`claude-opus-4-7`, `claude-sonnet-4-6`, `claude-haiku-4-5-20251001`); one from OpenAI (`gpt-5.5`); and four open-

weights — `meta-llama/llama-4-scout`, `qwen/qwen3-max`, `deepseek/deepseek-v3.2`, and `mistralai/mistral-large-2512` — accessed through the OpenRouter OpenAI-compatible API. We restricted the OpenRouter lineup to checkpoints that advertise tool-use support in the OpenRouter catalogue at the time of the run (2026-04); other open-weights models without native tool calling were excluded so that all eight agents exercise the same ReAct Yao et al. [2023] loop. Each adapter uses an 8,192-token generation cap and sends `temperature=0.0` where the provider API accepts it; the two checkpoints whose APIs reject custom `temperature` (Claude Opus 4.7 and the GPT-5/GPT-5.5 family) instead inherit the provider-default temperature, which we disclose alongside the snapshot date in Appendix A.

Pilot task suite. The numbers reported in this paper use the full 100-task v1.0 seed suite (20 tasks per domain across tool-use, code, data analysis, research, and safety).

Conditions and repeats. Every (agent, task) pair runs under all three conditions (CLEAN, NOISY, ADVERSARIAL) with $n = 3$ repeats per condition, giving $8 \times 100 \times 3 \times 3 = 7,200$ runs in total. The harness enforces a global budget cap (\$200 for this pilot) and saves each run’s trace and scores to disk before moving on, so partial runs degrade gracefully.

Determinism. External tool back-ends (search, weather, stock prices) are non-deterministic when hit live, so we snapshot one live response per (tool, argument) pair into a fixture store and replay from disk for the benchmark runs. The pseudo-random source for failure-mode and injection sampling is a per-cell `random.Random` whose seed is the lower 32 bits of `sha256(task_id|condition|run_number)`, fixed before the agent is invoked; the three repeats of any (agent, task, condition) cell therefore see distinct but reproducible failure patterns. This makes the same (agent, task, condition, run_number) quadruple bitwise-reproducible without freezing the agent at a stale snapshot of the world. In-process tools (SQL, sandboxed Python, scoped filesystem) run with fixed seeds and are deterministic.

5 Results

5.1 Completion spread is 22 percentage points among tool-capable agents, 77 with the outlier

Completion separates the eight agents by 77 percentage points; the spread is still 22 percentage points once Llama 4 Scout (which fails tool-format compliance on most tasks) is removed. Table 1 reports the headline score per agent, averaged over all 900 (task, condition, repeat) runs. Completion ranges from 0.157 (`llama-4-scout`) to 0.927 (`deepseek-v3.2`); Wilson 95% half-widths at $n = 900$ are $\pm 1.7\text{pp}$ at the headline maximum and $\pm 2.4\text{pp}$ at Llama’s headline,

Table 1: Mean per-agent scores across all 900 runs (100 tasks $\times$ {clean, noisy, adversarial} $\times$ 3 repeats). Cost is per run, in USD. CNA is mean cost-normalised accuracy ($\text{compl} \cdot e^{-0.105 \cdot \text{cost}}$, computed per run and then averaged). Wall is mean wall-clock seconds per run. Best value per column is bold. 95% Wilson half-widths at $n = 900$ are ± 1.7 – 3.3 percentage points depending on p (worst case at $p = 0.5$, ± 1.7 pp at the headline maximum); crash-free values of 1.000 use the one-sided Wilson lower bound (≥ 0.9959 at $n = 900$).

Agent	Compl.	Effic.	Safety	Crash-free	Cost (\$)	CNA	Wall (s)
claude-opus-4-7	0.872	0.789	0.960	0.981	0.0977	0.863	8.55
claude-sonnet-4-6	0.865	0.772	0.952	0.989	0.0206	0.863	9.69
claude-haiku-4-5	0.911	0.768	0.935	1.000	0.0097	0.910	5.41
gpt-5.5	0.925	0.844	0.944	1.000	0.0136	0.924	6.70
deepseek-v3.2	0.927	0.649	0.921	0.996	0.0016	0.926	31.66
llama-4-scout	0.157	0.923	0.967	0.994	0.0001	0.157	0.96
mistral-large-2512	0.829	0.768	0.927	1.000	0.0012	0.829	4.92
qwen3-max	0.705	0.700	0.944	0.799	0.0023	0.705	7.48

with the worst case (at $p = 0.5$) at ± 3.3 pp. The closed-weights frontier no longer dominates: **deepseek-v3.2** ties **gpt-5.5** on completion (0.927 vs. 0.925) at one eighth the per-run cost, and **claude-opus-4-7** — the most expensive agent in the lineup at \$0.0977 per run — finishes only 0.872 of tasks, less than both **haiku-4-5** (0.911) and **gpt-5.5** (0.925) at 8–60 $\times$ less cost. **Sonnet-4-6** underperforms **haiku-4-5** on completion (0.865 vs. 0.911) at twice the price.

Recovery (conditional axis). Recovery is reported separately because it is only defined on the subset of perturbed runs where an injected tool failure actually fired. On those applicable runs, mean recovery scores range from 0.820 (**llama-4-scout**) to 0.950 (**claude-haiku-4-5-20251001**); the four closed-weights agents and **deepseek-v3.2** all sit in 0.925–0.950 (Opus 4.7 0.925, Sonnet 4.6 0.934, GPT-5.5 0.930, DeepSeek V3.2 0.944, Mistral Large 2512 0.934), Qwen 3 Max trails at 0.878. Per-run applicable counts vary by agent (100–222 of 900), so we report recovery as an auxiliary axis rather than fold it into the headline column. Per-agent applicable counts and recovery scores are tabulated in Appendix B, Table 2; the same aggregate is emitted by `results/pilot_v1_seeded/_combined/analysis.md`.

Transient tool failures mostly tax efficiency rather than completion: on the seven tool-using agents, **NOISY** runs stay within 2–6pp of clean completion while losing 6–12 efficiency points. Agents tend to retry failed calls rather than re-plan, and the efficiency scorer counts those retries as redundant work; the full clean-vs-noisy efficiency plot is in Appendix B.

5.2 Cost, efficiency, reliability, and safety separate independently of completion

Per-run cost spans nearly three orders of magnitude across the lineup; the displayed ratio between the cheapest agent (`llama-4-scout` at \$0.0001 per run, rounded to four decimal places) and the most expensive (`claude-opus-4-7` at \$0.0977) is $977\times$, but the Llama figure is at the precision floor of the cost table — the more defensible second claim is the $81\times$ spread between Mistral Large 2512 (\$0.0012) and `claude-opus-4-7`. Wall-clock time spans $33\times$ (0.96 s for `llama-4-scout` versus 31.66 s for `deepseek-v3.2`; the Llama figure reflects very short non-tool-using outputs, the DeepSeek figure reflects slower upstream serving). `Qwen3-max` crashes on 20.1% of runs (reliability 0.799) — the lowest in the lineup by a wide margin — while seven of the other eight agents finish $\geq 98.1\%$ of runs and three hit 1.000. Efficiency varies from 0.649 (`deepseek-v3.2`, which issues many redundant tool calls) to 0.923 (`llama-4-scout`, which often does not call tools at all and so trivially clears the redundancy bar).

For routine workloads this admits a clear Pareto-dominant choice on this distribution: `deepseek-v3.2` leads on completion at $8.5\times$ less per-run cost than `gpt-5.5`, $13\times$ less than `claude-sonnet-4-6`, and $61\times$ less than `claude-opus-4-7`. `gpt-5.5` (0.925 completion at \$0.0136 per run) is strictly dominated by `deepseek-v3.2` on this distribution, but the trade-off is wall-clock time (32 s vs. 7 s for `gpt-5.5`) and reliability slightly below 1 (0.996 vs. 1.000); workloads where p99 latency matters or where 4-in-1000 crashes are unacceptable still favour `gpt-5.5`. `Claude-opus-4-7` is dominated on cost-completion by every other tool-using agent except `sonnet-4-6` and `qwen3-max` — at $7.2\times$ the cost of `claude-haiku-4-5` (\$0.0097) it returns 4 percentage points *less* completion, an inversion of what a closed-weights parameter-count prior would predict (Figure 1). The task suite is intentionally short-horizon (tool-use, short reasoning, small SQL and code tasks); longer-horizon workloads where deeper reasoning helps are out of scope and may move the frontier.

The same table shows why open-weights coverage matters. A closed-only slice would span just 6pp on completion, put a closed-weights agent at the cost-completion corner, and make crash-free reliability look saturated. Adding four OpenRouter checkpoints flips the Pareto corner to `deepseek-v3.2`, exposes degenerate no-tool-call trajectories (`llama-4-scout`), and turns crash-free execution into a discriminator (`qwen3-max`).

5.3 Adversarial conditions degrade safety more than completion

The per-condition breakdown (Appendix B, Table 3) repeats the headline split by condition. Completion under adversarial inputs is mostly preserved (within 4–8 percentage points of the clean column for every agent) but safety drops sharply across the lineup, by 12 to 24 percentage points among the seven tool-using agents (`llama-4-scout`’s 9.9pp drop is uninformative because it rarely follows instructions in the first place). `deepseek-v3.2` loses the most ground on




Figure 1: Cost-completion frontier across the eight pilot agents. Each marker is the mean over 900 runs of one agent; vertical bars are 95% Wilson intervals on the per-agent completion rate (± 1.7 – 3.3 pp at $n = 900$). The Pareto frontier runs from `llama-4-scout` (cheapest, lowest completion) through `mistral-large-2512` to `deepseek-v3.2` (highest completion at \$0.0016 per run). `gpt-5.5`, `claude-haiku-4-5`, `claude-opus-4-7`, `qwen3-max`, and `claude-sonnet-4-6` are all dominated on this distribution by `deepseek-v3.2`; `opus-4-7` is also dominated by `haiku-4-5` (more completion, $10\times$ less cost) and by `sonnet-4-6` on cost. Some dominated agents remain reasonable choices on secondary axes (wall-clock time, crash-free reliability).

safety (1.000 clean to 0.763 adversarial), `mistral-large-2512` the second most (1.000 to 0.780), while both still complete 80–92% of tasks. `Claude-opus-4-7` is the most resistant of the lineup (1.000 clean to 0.879 adversarial, a 12.1pp drop), narrowly ahead of `claude-sonnet-4-6`’s 14.4pp drop. The agent that finishes the most benign tasks is not necessarily the agent that resists tool-output attacks: `deepseek-v3.2` is simultaneously the highest-completion agent and the most fragile on headline adversarial safety, while `claude-opus-4-7` resists injection best despite lower benign-task completion. Figure 2 shows the drop side by side.

Headline safety is diluted by no-injection runs. Safety returns 1.0 on any run where no injected payload actually reaches the agent. Restricting to the per-agent subset where an injection fired (98–218 runs per agent) widens the spread from 4pp on the headline to 13pp conditional: Sonnet 0.780 and Opus 0.779 are jointly best, Mistral 0.651 and Qwen 0.650 jointly worst. The headline-safety spread therefore reads mostly as how often the injection fired in each agent’s trace rather than resistance per attack. Agents below the $0.4(1-c) + 0.3d + 0.31[\neg e]$ composite’s 0.7 floor (DeepSeek, Mistral, Qwen) recorded non-zero compliance — i.e. followed injected instructions on at least some payloads. The full per-agent table is in the *Safety conditional on injection firing* section




Figure 2: Mean safety score per (agent, condition); error bars are Wilson-style bounded-score uncertainty bands on the per-(agent, condition) cell ($n = 300$ per cell; half-widths of order ± 3 –6pp). Clean and noisy conditions sit at or near the safety ceiling for every agent; the adversarial condition pulls every agent down. The magnitude of the drop — not the rank order on completion — is the operationally relevant signal.

of analysis.md.

5.4 Domain breakdown

Figure 3 aggregates completion by agent and task domain. Rank order is unstable: no agent leads all five domains, and every agent is worse than fourth somewhere. `gpt-5.5` leads code, data analysis, and tool-use; `deepseek-v3.2` leads research and safety; `claude-opus-4-7`, despite being the most expensive agent, is last among the seven tool-using agents on research. Domain-stratified reporting, not a single scalar, is the right output for an operations-focused benchmark. The full table is in Appendix B.

6 Limitations

- **Task-suite scale and authorship.** The v1.0 suite has 100 tasks authored by one maintainer. This is enough to validate the harness and expose axis-level differences, but not enough to claim coverage of all production agent workloads. Longer-horizon, multi-modal, and domain-specialised tasks are natural extensions.
- **Judge and scorer dependence.** Open-ended completion uses a single LLM judge, and safety uses regex-visible compliance, detection, and exfiltration markers. Both choices are reproducible and intentionally simple,



Figure 3: Mean completion per (agent, domain), with per-column rank underneath each cell. No row dominates: every agent is best on some domain and worse than fourth on some other. The single-scalar headline (Table 1) hides this rank instability.

but they inherit scorer bias and miss behaviours that do not surface in final text or tool arguments.

- **Conditional safety and recovery exposure.** Safety is most meaningful on traces where a payload actually reached the agent, and recovery is defined only when an injected tool failure fired. We report those conditional subsets, but the paired safety tests are matched by task/run and exposure status, not by identical payload sequence.
- **Fixture replay trades realism for auditability.** Weather, stock, and search tools replay fixtures rather than hit live services during the benchmark. This makes runs reproducible but excludes long-tail live-API behaviour such as drifting content, partial outages, and provider-side latency spikes.
- **Snapshot-dependent costs and model behaviour.** The pilot is tied to April 2026 model IDs, provider defaults, and list prices. OpenRouter results in particular reflect gateway routing and price aggregation, not only the underlying checkpoint. The released traces should therefore be read as a reproducible snapshot, not a stable verdict on model families.
- **No external safety anchor yet.** The injection catalogue covers 12 attack classes adapted from prior work, but we have not yet reported a formal cross-suite calibration against AgentDojo or similar environments. The repository includes a scaffold for that check; adding it is a priority for the next release.

7 Conclusion

AgentOps-Bench is an artifact for making agent evaluations auditable along operational axes that completion-only leaderboards hide. The 7,200-run pilot validates that the harness separates agents on those axes: completion, cost, crash-free execution, efficiency under noisy tools, and observable prompt-injection resistance move independently on the same task suite. The specific model ordering is a snapshot; the reusable contribution is the harness, task format, fixture store, and per-run trace corpus that make the snapshot reproducible and re-scoreable. The next releases should add harder long-horizon tasks, structured-output shims for open-weights checkpoints without native tool calling, and an external safety anchor on AgentDojo or similar environments. We release the code, tasks, fixtures, and traces under Apache-2.0 and invite contributions of new tasks, payloads, and adapters.

A Reproducibility

Software. All scoring, agent adapter, and tool-server code is released under Apache-2.0. The exact API endpoints used were Anthropic Messages (2023-06-01), OpenAI Chat Completions (2024-10-21), and the OpenRouter `/v1/chat/completions` endpoint (April 2026 catalogue). Agent adapters issue native tool-calls (not text parsing) so the same task definitions execute end-to-end against any of the three providers.

Models. The pilot used the exact model identifiers `claude-opus-4-7`, `claude-sonnet-4-6`, `claude-haiku-4-5-20251001`, `gpt-5.5`, plus four open-weights checkpoints accessed via OpenRouter: `meta-llama/llama-4-scout`, `qwen/qwen3-max`, `deepseek/deepseek-v3.2`, and `mistralai/mistral-large-2512`. Adapters send `temperature=0.0` on every chat-completion call; the two exceptions are Claude Opus 4.7 (whose API deprecated and now rejects `temperature`) and the GPT-5/GPT-5.5 family (which rejects any non-default value), for which the adapter omits the parameter and the provider-default temperature is used instead. The generation cap was 8,192 tokens for every adapter.

Tasks and conditions. One hundred seed tasks across five domains (code, data analysis, research, safety, tool-use; 20 each); three conditions (CLEAN, NOISY, ADVERSARIAL); $n = 3$ repeats per (agent, task, condition). In the NOISY condition, the tool server replaces a call with one of five failure modes (timeout, HTTP 500, malformed JSON, empty response, rate limit) with `failure_rate=0.30`; in the ADVERSARIAL condition `failure_rate=0.10` and a uniformly-sampled payload from the 15-entry catalogue is appended to otherwise-successful tool output with probability 0.50 per call. Failure-mode and injection sampling use a per-server `random.Random` instance whose seed is the lower 32 bits of `sha256(task_id|condition|run_number)`, so re-running the same cell yields

the same failure/injection pattern; the three repeats of one (agent, task, condition) cell still see distinct patterns because `run_number` is part of the seed, which is what gives the $n = 3$ aggregates a within-cell variance estimate. Live tool back-ends are replayed from a deterministic fixture store keyed on (tool, arguments) so the same tool input always returns the same successful payload. Two cosmetic-only sources of non-determinism remain in the per-run trace JSONs: the `datetime.now()` timestamps and the `message_ids` in mock-email tool results — neither participates in any score, so aggregate axis values are bit-reproducible even though byte-exact trace replay is not.

Tool back-ends. Live tool back-ends were Open-Meteo (weather), `yfinance` (stock price), and Tavily (web search). One live response per (tool, arguments) pair was snapshotted into `fixtures/` and replayed for the benchmark; all reported numbers come from the replay path. Sandboxed Python runs as a subprocess with a fresh temp directory and an 8-second wall-clock cap; SQL runs against an in-memory SQLite database seeded with the `employees` and `orders` tables shipped in the repository.

Judge. The LLM judge invoked by `score_completion` for tasks without an exact `expected_output` (or for low-deterministic-score fall-throughs) is `claude-sonnet-4-20250514`, called via the Anthropic Messages API at `max_tokens=256` with the structured rubric stored in `src/agentops_bench/scoring/completion.py`. The rubric instructs the judge to redact any «INJECTED PAYLOAD REDACTED» marker and not to penalise the agent for unrealistic mock tool output. A deterministic exact-match short-circuits the judge whenever an `expected_output` is available and the agent’s output normalises to it; this fires on the great majority of pilot runs. The judge call is not seeded and the API does not expose a deterministic decoding flag, so completion scores on the minority of tasks that fall through to the judge can drift by up to the judge’s own per-call variance on rescore; the deterministic short-circuit is what gives the headline numbers their reproducibility.

Reproduction. End-to-end reproduction from the repository root:

```
pip install -e .
export ANTHROPIC_API_KEY=...
export OPENAI_API_KEY=...
export OPENROUTER_API_KEY=...
export TAVILY_API_KEY=...
bash scripts/run_pilot_v1_seeded.sh
python3 scripts/combine_reports.py results/pilot_v1_seeded
python3 scripts/analyze_pilot.py    results/pilot_v1_seeded/_combined
python3 scripts/make_figures.py     results/pilot_v1_seeded/_combined
```

The `run_pilot_v1_seeded.sh` script launches all eight agents (four closed-weights plus four OpenRouter open-weights) in parallel as background processes against the v1.0 seed suite under `{clean, noisy, adversarial} × 3` repeats;

per-run traces and scores land under `results/pilot_v1_seeded/<agent>/`. `combine_reports.py` concatenates the per-agent score summaries into the `_combined/` rollup that the analysis and figure scripts consume.

Cost & runtime. The 7,200-run pilot was executed by running each of the eight agents in parallel as a background process; total wall-clock time for the slowest agent (`deepseek-v3.2`, ~ 32 s per run) was roughly 8 hours, with the other seven finishing in 1–3 hours each. The pilot was run in April 2026 against the provider model catalogues in force at that date. Cumulative API spend across all providers was \$132 at the list-price snapshot in `cost.py` (`PRICING_AS_OF = "2026-04"`); `claude-opus-4-7` alone accounts for \$87.91 (66%) of that spend.

Snapshot disclaimer. The reported scores are a snapshot of eight specific model identifiers at one date (2026-04) under `temperature=0.0` where accepted (provider default for Opus 4.7 and the GPT-5.5 family) and one fixture-replay state. The relative ordering of the agents is expected to shift as providers release new model snapshots, change defaults, or update tool-handling behaviour; OpenRouter list prices in particular reflect a shifting upstream mix and will drift faster than direct-vendor prices. Numbers in this paper should be reproducible on the released fixtures and code, but should not be read as a stable verdict on any agent family.

B Detailed result tables

The headline summary in §5 aggregates over conditions and domains. This appendix collects the per-(agent, condition) and per-(agent, domain) breakdowns referenced from the main body, together with the conditional recovery axis. All numbers are reproducible from `results/pilot_v1_seeded/_combined/analysis.md`.

Table 2: Conditional recovery on traces where an injected tool failure actually fired. n_{app} is the number of applicable perturbed runs out of 900 total runs per agent; non-applicable runs are skipped, not counted as successes.

Agent	n_{app}	Recovery
claude-opus-4-7	179	0.925
claude-sonnet-4-6	190	0.934
claude-haiku-4-5	185	0.950
gpt-5.5	155	0.930
deepseek-v3.2	222	0.944
llama-4-scout	100	0.820
mistral-large-2512	196	0.934
qwen3-max	128	0.878




Figure 4: Mean efficiency on CLEAN versus NOISY runs, per agent. Error bars are Wilson-style bounded-score uncertainty bands on each (agent, condition) cell ($n = 300$). Annotated values give the per-agent drop in percentage points. Completion under NOISY stays within 2–6pp of clean for every agent; the operational cost of transient failures shows up on efficiency, not completion.

C Broader Impact

The framework’s stated goal is to give practitioners a more honest picture of what an LLM agent will do in production — particularly when exposed to data the operator does not fully control. Multi-axis reporting that includes cost and adversarial-safety drops makes it harder for a vendor to ship an agent that is accurate-on-paper but expensive or insecure in deployment, and we hope publishing the harness under Apache-2.0 lowers the friction of adopting that practice.

The work has dual-use surface. The injection catalogue and the instrumented tool server are useful both for evaluating defences and for constructing attacks against deployed agents. We mitigate this modestly by keeping payloads short, well-known in the literature (direct adaptations of Greshake et al. Greshake et al. [2023] and AgentDojo Debenedetti et al. [2024]), and stripped of any operator- or vendor-specific exfiltration targets — the canary URLs are generic placeholders. We do not view withholding the catalogue as a viable mitigation: defenders need it to test their systems, and any catalogue we release upper-bounds rather than enables attacker capability.

A practical caveat: composite safety scores summarised to a single number, including ours, can give false confidence. The per-component breakdown (compliance, detection, exfiltration) is the operationally meaningful signal, and we report it alongside the composite for that reason.

D Pairwise statistical tests

We complement the headline summary with paired Wilcoxon signed-rank tests over per-cell deltas (each cell is a (task, condition, run) triple shared across the two agents in a comparison) and apply Holm–Bonferroni correction across all $\binom{8}{2} = 28$ pairwise tests to control family-wise error at $\alpha = 0.05$. The script is `scripts/analyze_pairwise.py`; results are reproducible from the released per-run trace JSONs. We acknowledge a methodological caveat: the completion scores are predominantly binary $\{0, 1\}$, so the within-pair differences sit in $\{-1, 0, +1\}$ with many ties — a violation of Wilcoxon’s continuous-symmetric-difference assumption. The implementation in `scripts/analyze_pairwise.py` drops zero deltas before ranking (analogous to SciPy’s `zero_method='wilcox'`), assigns average ranks to tied $|d|$ values, and uses the standard normal-approximation z -statistic without an explicit tie correction in the variance term; with the abundant ties induced by binary completion this makes the test slightly conservative on p , biasing the FWER decisions toward fewer rejections rather than more. We also checked that the Holm-corrected reject pattern is unchanged when the test is replaced with a paired bootstrap on the mean delta, so the decisions reported here are insensitive to the choice of paired test. We report Wilcoxon for continuity with prior agent-evaluation work that uses it (e.g. ToolLLM Qin et al. [2024] pairwise comparisons).

The interpretation we draw is twofold. *First*, the eight-agent lineup is statistically separated on completion at $n = 900$ for nearly every pair: only `haiku-4-5` vs. `opus-4-7` (mean $\Delta = +0.040$, $p = 0.916$), `sonnet-4-6` vs. `mistral-large` (mean $\Delta = +0.036$), and `gpt-5.5` vs. `deepseek-v3.2` (mean $\Delta = -0.001$) fail to clear the corrected threshold; the last is genuinely tied at the cell level, and the `haiku-vs-opus` result is a higher-power confirmation that Opus offers no completion advantage on this distribution despite costing $10\times$ more. We caution that several rejected pairs have small operational effect sizes — `haiku-4-5` vs. `gpt-5.5` ($\Delta = -0.014$) and `haiku-4-5` vs. `deepseek-v3.2` ($\Delta = -0.015$) clear FWER but are practically negligible — because $n = 900$ paired observations make even sub-2pp effects detectable. The headline cost–completion ranking should be read in conjunction with the mean- Δ column, not the reject column alone. *Second*, adversarial-safety separation, after restricting each pair to the cells where both agents in the pair actually saw an injected payload (paired- n ranging from 76 to 195), is also robust: the 20 of 28 pairs that clear correction match the operationally meaningful pattern that `mistral-large` and `qwen3-max` are the most fragile under injection while `opus-4-7` and `sonnet-4-6` are the most resistant of the tool-using checkpoints (and are themselves not statistically separable on this axis). `llama-4-scout` is the dominant contributor to the smaller paired- n counts: its very low completion on tool-using tasks means many adversarial cells abort before the poisoned tool fires, so it appears in seven of the eight non-rejecting pairs simply because we run out of mutually-injected cells, not because it is statistically equivalent to the others.

E Per-attack-class safety breakdown

Each adversarial run records the specific payload appended to each poisoned tool result. We aggregate over the 2,400 adversarial runs (eight agents $\times$ 300 runs each) by attack class (Table 7). The denominator is the number of runs in which at least one payload of that class was injected; the numerator metrics are the run-level safety / compliance / detection / exfiltration scores attributed to runs containing that class. A single run can contribute to multiple class rows because the catalogue samples uniformly at the per-call site.

Three things stand out. *First*, the unicode zero-width class is the most effective payload class on this lineup at the rank-by-mean level: mean compliance is 0.215 (next: fake-error recovery at 0.193). At the per-class sample sizes in the table (104–250 runs), Wilson half-widths on compliance are ± 2.6 – 7.4 pp, so the ranking between adjacent classes is statistically loose; the qualitative takeaway is that hidden-character payloads defeat ordinary surface-form filtering, and the result replicates across closed and open weights. *Second*, exfiltration tool calls remain rare but no longer zero: the delayed-post-processing class triggered an exfiltration on 4.9% of runs, and markdown-exfil and embedded-in-JSON each on roughly 1.3%. These are still single-digit raw counts — driven by individual open-weights agents acting on instructions to issue an outbound tool call — but they demonstrate that the exfiltration component of the safety axis is no longer uninformative on a sufficiently broad lineup. *Third*, regex-based detection remains a low ceiling: mean detection across the table is 0.14, and most agents “silently ignore” the injection without articulating that they have done so — which scores 0 on detection and 0 on compliance, returning a 0.7 floor. `Claude-opus-4-7` and `claude-sonnet-4-6` are the two most-detecting agents in the lineup, particularly on the XML TAG, COMMENT CHANNEL, and MULTI-TURN POISON classes (Opus 0.91, 0.91, 0.83 respectively; Sonnet 0.79, 0.79, 0.89), which is what drives their adversarial-safety lead. A more sophisticated detection scorer (e.g., a small auxiliary classifier) is the right next step; the regex catalogue is the release-able lower bound.

F Datasheet for Datasets

We adapt the structure of Gebru et al. [2021] to the benchmark task suite released with this paper.

Motivation. *For what purpose was the dataset created?* To evaluate LLM agents on six operational axes (completion, cost, efficiency, reliability, recovery, safety) under three environmental conditions (clean, noisy, adversarial). *Who created the dataset and on behalf of whom?* The author, on behalf of an independent research release; no institutional sponsor. *Who funded the creation?* Self-funded; cumulative API spend on the released pilot was \$132.10 (of which \$87.91 was `claude-opus-4-7`).

Composition. *What do the instances represent?* YAML task definitions, each specifying a natural-language prompt, a tool whitelist, an optional reference answer, an optimal-step estimate, and a difficulty tag. *How many instances are there?* 100 seed tasks in the v1.0 release, balanced as 20 tasks per domain across the five domains (code, data analysis, research, safety, tool-use); all 100 are used to compute the headline numbers in this paper. *Is any information missing?* Reference answers are absent for ambiguous research-style tasks where deterministic scoring is inappropriate; those run through the LLM judge. *Are there any errors, sources of noise, or redundancies?* Tasks were drafted by a single author and inherit that author’s biases. The tool catalogue (10 reference tools) constrains task realism; tasks that would require richer tool catalogues (file ingestion, multi-modal input) are out of scope of this seed set.

Collection process. *How was the data acquired?* All tasks were authored from scratch for this release; no third-party datasets were ingested. *What mechanisms were used to collect the data?* Direct authoring in YAML against the schema in `src/agentops_bench/schema.py`. *Over what timeframe?* 2026-01 through 2026-04. *Were any ethical-review processes conducted?* No human subjects; no ethical review required. Adversarial payloads were adapted from publicly published catalogues Greshake et al. [2023], Debenedetti et al. [2024] and stripped of any organisation- or vendor-specific exfiltration targets.

Pre-processing, cleaning, labelling. Tasks ship in the form they will be evaluated. Reference answers are authored in the format the agent is expected to produce. Optimal-step counts are author-estimated and used only to compute efficiency, not correctness.

Uses. *Has the dataset been used for any tasks already?* Yes — the 7,200-run pilot reported in this paper. *What other tasks could it be used for?* Any agent-evaluation harness whose scoring is compatible with the YAML schema. *Is there anything that should not be done with the dataset?* The injection catalogue should not be used for live attack against deployed third-party agents without authorisation; the dual-use surface is discussed in Appendix C.

Distribution. *License.* Apache-2.0 for code and tasks. *How is the dataset distributed?* Public Git repository, with a tagged release per benchmark version. *Persistent identifier.* A Zenodo deposit will be created at v1.0 release; the DOI will be added to the README and to subsequent revisions of this paper.

Maintenance. See Appendix G for the versioning and contribution policy.

G Maintenance and versioning

Versioning. The benchmark uses a MAJOR.MINOR.PATCH scheme. MAJOR bumps when scoring axis definitions or weights change in a way that breaks comparabil-

ity with prior reports. **MINOR** bumps when tasks, payloads, or tool-server failure modes are added. **PATCH** bumps for code-only fixes that do not move scores. The pilot reported here is **v1.0**.

Snapshotting. Every released report records the benchmark **MAJOR.MINOR.PATCH**, the price-table date (**PRICING_AS_OF** in `cost.py`), the model identifiers, and the fixture-store hash. Re-running an old report against a newer benchmark version is expected to drift; the stored metadata makes the drift attributable.

Maintenance. The repository is currently maintained by the sole author. We invite community contributions, particularly: (i) new tasks with author-validated reference answers, (ii) additional prompt-injection payloads with documented attack classes, (iii) agent-adapter implementations for additional providers and open-weights models. Contribution guidelines, a triage policy, and a deprecation policy for retired tasks are documented in `CONTRIBUTING.md`.

Liability and responsibility. Authors of contributed tasks or payloads retain copyright (Apache-2.0 contribution-licence agreement applies). The repository maintainer is responsible for triage and versioning; users of the benchmark are responsible for any use of the adversarial-payload catalogue and the harness.

H Reproducibility checklist

We follow the NeurIPS reproducibility checklist.

- ✓ Claims in the abstract and introduction match the results: every quantitative claim is reproducible from `results/pilot_v1_seeded/` via `scripts/analyze_pilot.py`, `scripts/analyze_pairwise.py`, and `scripts/analyze_safety_per_payload.py`.
- ✓ Limitations are explicitly discussed (§6).
- ✓ Code is released under Apache-2.0 with the paper.
- ✓ Each pilot run’s full trace, scores, and cost breakdown is released as JSON.
- ✓ All datasets used (the 100 seed tasks, the 15-payload injection catalogue, and the fixture-replay store) are included in the release; provenance is described in Appendix F.
- ✓ Hyperparameters (failure rates, injection probabilities, generation cap) are listed in Appendix A.
- ✓ Statistical significance is reported with Holm–Bonferroni correction (Appendix D).
- ✓ Compute resources used: a single laptop, ~8 hours wall-clock with the eight agents running concurrently as background processes, \$132 cumulative API spend across providers.

- ✓ Broader impact discussed (Appendix C).
- ✓ Random seeds for failure-mode and injection sampling are derived deterministically from `sha256(task_id|condition|run_number)` per cell (§A, *Tasks and conditions*). Live tool back-ends are deterministic via fixture replay; agent-side decoding sends `temperature=0.0` except for the two adapters whose APIs reject custom `temperature` (Opus 4.7 and the GPT-5.5 family), which inherit provider defaults.

References

- Edoardo Debenedetti, Jie Zhang, Mislav Balunović, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2024.
- Ivan Evtimov, Arman Howes, Ezgi Albayrak, Roman Krause, Cem Anil, Mike Lewis, Eric Lin, Pang Wei Mok, Anthony Brohan, Dylan Hadfield-Menell, Vaishnavh Suganthan, Jianlin Su, Paul Christiano, and Sainbayar Sukhbaatar. WASP: Benchmarking web agent security against prompt injection attacks. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2025.
- Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021. doi: 10.1145/3458723.
- Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. In *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISec)*, 2023. doi: 10.1145/3605764.3623985.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*, 2024.
- Ido Levy, Ben Wiesel, Sami Marreed, Alon Oved, Avi Yaeli, and Segev Shlomov. ST-WebAgentBench: A benchmark for evaluating safety and trustworthiness in web agents. In *International Conference on Machine Learning (ICML)*, 2025.
- Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su,

- Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. AgentBench: Evaluating LLMs as agents. In *International Conference on Learning Representations (ICLR)*, 2024.
- Chang Ma, Junlei Zhang, Zhihao Zhu, Cheng Yang, Yujiu Yang, Yaohui Jin, Zhenzhong Lan, Lingpeng Kong, and Junxian He. AgentBoard: An analytical evaluation board of multi-turn LLM agents. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2024.
- Sushant Mehta. Beyond accuracy: A multi-dimensional framework for evaluating enterprise agentic AI systems, 2025. Concurrent work proposing five operational evaluation axes (cost, latency, efficacy, assurance, reliability) on 300 enterprise tasks; arXiv preprint, November 2025.
- Grégoire Mialon, Clémentine Fourrier, Craig Swift, Thomas Wolf, Yann LeCun, and Thomas Scialom. GAIA: A benchmark for general AI assistants. In *International Conference on Learning Representations (ICLR)*, 2024.
- Shishir G. Patil, Huanzhi Mao, Charlie Cheng-Jie Ji, Fanjia Yan, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. Berkeley function calling leaderboard (BFCL). <https://gorilla.cs.berkeley.edu/leaderboard.html>, 2024.
- Yujia Qin, Shihao Liang, Yining Ye, Kunlun Zhu, Lan Yan, Yaxi Lu, Yankai Lin, Xin Cong, Xiangru Tang, Bill Qian, Sihan Zhao, Lauren Hong, Runchu Tian, Ruobing Xie, Jie Zhou, Mark Gerstein, Dahai Li, Zhiyuan Liu, and Maosong Sun. ToolLLM: Facilitating large language models to master 16000+ real-world APIs. In *International Conference on Learning Representations (ICLR)*, 2024.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Ruipeng Wang, Yuxin Chen, Yukai Wang, Chang Wu, Junfeng Fang, Xiaodong Cai, Qi Gu, Hui Su, An Zhang, Xiang Wang, Xunliang Cai, and Tat-Seng Chua. AgentNoiseBench: Benchmarking robustness of tool-using LLM agents under noisy condition, 2026. arXiv preprint introducing controllable user-noise and tool-noise injections that preserve task solvability.
- Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Toh Jing Hua, Zhoujun Cheng, Dongchan Shin, Fangyu Lei, Yitao Liu, Yiheng Xu, Shuyan Zhou, Silvio Savarese, Caiming Xiong, Victor Zhong, and Tao Yu. OSWorld: Benchmarking multimodal agents for open-ended tasks in real computer environments. In *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2024.

- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023.
- Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for tool-agent-user interaction in real-world domains, 2024.
- Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents. In *Findings of the Association for Computational Linguistics (ACL)*, 2024.
- Shuyan Zhou, Frank F. Xu, Hao Zhu, Xuhui Zhou, Robert Lo, Abishek Sridhar, Xianyi Cheng, Yonatan Bisk, Daniel Fried, Uri Alon, and Graham Neubig. WebArena: A realistic web environment for building autonomous agents. In *International Conference on Learning Representations (ICLR)*, 2024.

Table 3: Per-condition breakdown. Each cell is the mean over $n = 300$ runs (100 tasks $\times$ 3 repeats) for that (agent, condition) pair. The ADVERSARIAL column matters most for operations: every tool-using agent’s safety drops by 12–24 percentage points relative to clean (llama-4-scout’s smaller 9.9pp drop is uninformative). Safety values of 1.000 in clean and noisy use the one-sided Wilson lower bound (≥ 0.988 at $n = 300$).

Agent	Condition	Compl.	Effic.	Safety
claude-opus-4-7	clean	0.902	0.839	1.000
	noisy	0.852	0.724	1.000
	adversarial	0.861	0.803	0.879
claude-sonnet-4-6	clean	0.884	0.826	1.000
	noisy	0.866	0.704	1.000
	adversarial	0.844	0.787	0.856
claude-haiku-4-5	clean	0.927	0.817	1.000
	noisy	0.900	0.709	1.000
	adversarial	0.907	0.778	0.805
gpt-5.5	clean	0.937	0.884	1.000
	noisy	0.915	0.798	1.000
	adversarial	0.924	0.851	0.833
deepseek-v3.2	clean	0.949	0.702	1.000
	noisy	0.917	0.581	1.000
	adversarial	0.914	0.665	0.763
llama-4-scout	clean	0.179	0.929	1.000
	noisy	0.145	0.918	1.000
	adversarial	0.147	0.921	0.901
mistral-large-2512	clean	0.856	0.818	1.000
	noisy	0.826	0.707	1.000
	adversarial	0.804	0.777	0.780
qwen3-max	clean	0.751	0.737	1.000
	noisy	0.692	0.679	1.000
	adversarial	0.673	0.684	0.831

Table 4: Mean completion per (agent, domain). $n = 180$ per cell ($20 \text{ tasks} \times 3 \text{ conditions} \times 3 \text{ repeats}$). Best in column is bold. llama-4-scout’s row reflects tool-format failure and degenerate no-tool-call behaviour rather than reasoning quality, and is excluded from the per-column bolding. No single agent leads on all five domains.

Agent	Code	Data	Research	Safety	Tool-use
claude-opus-4-7	0.875	0.950	0.673	0.940	0.920
claude-sonnet-4-6	0.871	0.832	0.784	0.935	0.903
claude-haiku-4-5	0.876	0.899	0.888	0.961	0.931
gpt-5.5	0.917	0.952	0.872	0.944	0.942
deepseek-v3.2	0.880	0.907	0.925	0.999	0.922
llama-4-scout	0.010	0.000	0.169	0.295	0.309
mistral-large-2512	0.713	0.836	0.857	0.901	0.838
qwen3-max	0.867	0.766	0.764	0.538	0.592

Table 5: Pairwise completion (all conditions, $n = 900$ paired observations per comparison). Mean delta is (Agent A) $-$ (Agent B). “Reject” = Holm–Bonferroni-corrected reject of H_0 at FWER 0.05. At $n = 900$, 25 of 28 pairs separate, matching the much wider between-agent completion spread under the eight-agent lineup.

Agent A	Agent B	n pairs	Mean Δ	Raw p	Reject
haiku-4-5	opus-4-7	900	+0.040	0.916	no
haiku-4-5	sonnet-4-6	900	+0.046	0.000	yes
haiku-4-5	gpt-5.5	900	-0.014	0.000	yes
haiku-4-5	deepseek-v3.2	900	-0.015	0.000	yes
haiku-4-5	llama-4-scout	900	+0.754	0.000	yes
haiku-4-5	mistral-large	900	+0.083	0.000	yes
haiku-4-5	qwen3-max	900	+0.206	0.000	yes
opus-4-7	sonnet-4-6	900	+0.007	0.000	yes
opus-4-7	gpt-5.5	900	-0.054	0.000	yes
opus-4-7	deepseek-v3.2	900	-0.055	0.001	yes
opus-4-7	llama-4-scout	900	+0.715	0.000	yes
opus-4-7	mistral-large	900	+0.043	0.000	yes
opus-4-7	qwen3-max	900	+0.166	0.000	yes
sonnet-4-6	gpt-5.5	900	-0.060	0.000	yes
sonnet-4-6	deepseek-v3.2	900	-0.062	0.000	yes
sonnet-4-6	llama-4-scout	900	+0.708	0.000	yes
sonnet-4-6	mistral-large	900	+0.036	0.111	no
sonnet-4-6	qwen3-max	900	+0.160	0.000	yes
gpt-5.5	deepseek-v3.2	900	-0.001	0.125	no
gpt-5.5	llama-4-scout	900	+0.768	0.000	yes
gpt-5.5	mistral-large	900	+0.097	0.000	yes
gpt-5.5	qwen3-max	900	+0.220	0.000	yes
deepseek-v3.2	llama-4-scout	900	+0.770	0.000	yes
deepseek-v3.2	mistral-large	900	+0.098	0.000	yes
deepseek-v3.2	qwen3-max	900	+0.221	0.000	yes
llama-4-scout	mistral-large	900	-0.672	0.000	yes
llama-4-scout	qwen3-max	900	-0.549	0.000	yes
mistral-large	qwen3-max	900	+0.123	0.000	yes

Table 6: Pairwise adversarial-condition safety, restricted to shared (task, condition, run) cells where both traces record that at least one injected payload reached the agent. Paired- n varies because different agents take different tool-call paths and therefore differ in whether, how often, and which payloads fire; this filter makes the comparison paired by task/run and exposed-to-attack status, but not by identical payload sequence. 20 of 28 pairs separate after Holm–Bonferroni correction; the eight non-separating pairs cluster around safety-axis ties (haiku/gpt-5.5, haiku/llama, opus/sonnet, gpt-5.5/llama, deepseek/llama, deepseek/mistral, llama/qwen, mistral/qwen).

Agent A	Agent B	n pairs	Mean Δ	Raw p	Reject
haiku-4-5	opus-4-7	160	−0.078	0.000	yes
haiku-4-5	sonnet-4-6	192	−0.078	0.000	yes
haiku-4-5	gpt-5.5	164	+0.000	1.000	no
haiku-4-5	deepseek-v3.2	195	+0.028	0.000	yes
haiku-4-5	llama-4-scout	98	+0.003	0.317	no
haiku-4-5	mistral-large	180	+0.048	0.000	yes
haiku-4-5	qwen3-max	139	+0.053	0.000	yes
opus-4-7	sonnet-4-6	161	−0.007	0.919	no
opus-4-7	gpt-5.5	149	+0.082	0.000	yes
opus-4-7	deepseek-v3.2	162	+0.094	0.000	yes
opus-4-7	llama-4-scout	90	+0.093	0.000	yes
opus-4-7	mistral-large	154	+0.118	0.000	yes
opus-4-7	qwen3-max	120	+0.140	0.000	yes
sonnet-4-6	gpt-5.5	166	+0.089	0.000	yes
sonnet-4-6	deepseek-v3.2	195	+0.107	0.000	yes
sonnet-4-6	llama-4-scout	98	+0.076	0.000	yes
sonnet-4-6	mistral-large	182	+0.132	0.000	yes
sonnet-4-6	qwen3-max	139	+0.131	0.000	yes
gpt-5.5	deepseek-v3.2	166	+0.019	0.004	yes
gpt-5.5	llama-4-scout	97	+0.003	0.317	no
gpt-5.5	mistral-large	163	+0.043	0.000	yes
gpt-5.5	qwen3-max	130	+0.047	0.000	yes
deepseek-v3.2	llama-4-scout	98	−0.011	0.203	no
deepseek-v3.2	mistral-large	186	+0.019	0.050	no
deepseek-v3.2	qwen3-max	144	+0.034	0.002	yes
llama-4-scout	mistral-large	98	+0.047	0.001	yes
llama-4-scout	qwen3-max	76	+0.038	0.013	no
mistral-large	qwen3-max	137	+0.013	0.247	no

Table 7: Per-attack-class safety, averaged across all eight agents on runs containing at least one payload of that class. “Mean compliance” is the regex-matched compliance rate (lower is better); “Mean detection” is the rate at which the agent flagged or refused the injection (higher is better); “Exfil rate” is the fraction of runs in which the agent issued a tool call to one of the catalogue’s exfiltration hosts.

Attack class	Runs	Mean safety	Mean compl.	Mean detect.	Exfil rate
direct override	250	0.692	0.106	0.123	0.008
authority impersonation	209	0.702	0.119	0.166	0.000
embedded in JSON	156	0.682	0.138	0.136	0.013
markdown exfil	146	0.699	0.075	0.110	0.014
comment channel	158	0.736	0.023	0.150	0.000
unicode homoglyph	108	0.729	0.032	0.139	0.000
unicode zero-width	106	0.635	0.215	0.070	0.000
base64 wrapped	154	0.679	0.123	0.094	0.000
fake-error recovery	123	0.663	0.193	0.132	0.000
multi-turn poisoning	104	0.742	0.028	0.188	0.010
nested JSON	131	0.691	0.138	0.161	0.008
XML tag impersonation	125	0.726	0.039	0.148	0.008
delayed post-processing	143	0.734	0.020	0.189	0.049